

Transparency in Neuro-Oncology: A Quantitative Framework for Explainable AI in Brain Tumor Classification

Shek Lun Leung
Quantum AI Laboratory
sheklunleung.qai@proton.me

Sai Oop Mong, MD
Department of Neuroradiology
saioomp Mong88@gmail.com

2026

Abstract

While Deep Convolutional Neural Networks (CNNs) have demonstrated dermatologist-level efficacy in medical image classification, their reliable deployment in neuro-oncology is impeded by a lack of interpretability. This study introduces a unified framework for multi-class brain tumor classification (Glioma, Meningioma, Pituitary, No Tumor) using T1-weighted MRI, integrated with a rigorous Explainable AI (XAI) verification suite. Unlike prior works that rely on qualitative visual inspection, we quantify the *faithfulness* and *robustness* of our explanations using occlusion sensitivity (Faithfulness Score: 0.37) and geometric perturbation analysis (Rotation Stability: 0.98). By upgrading from standard Grad-CAM to **Grad-CAM++**, we demonstrate statistically significant improvements in lesion localization. Our findings propose a validated pathway for bridging the gap between algorithmic probability and clinical trust.

1 Introduction

The rapid stratification of intracranial neoplasms is critical for optimizing patient outcomes. Structural Magnetic Resonance Imaging (MRI), specifically T1-weighted sequences, serves as the primary modality for anatomical assessment. However, the diagnostic workflow is labor-intensive and susceptible to inter-observer variability.

Although CNNs such as ResNet and DenseNet have achieved high diagnostic accuracy [?], they function as "black boxes." In high-stakes clinical environments, a prediction without rationale is actionable only if it can be verified [?]. A false positive triggered by a skull artifact or scanner noise is clinically unacceptable.

This paper addresses three critical gaps in the current literature:

1. **High-Fidelity Explanation:** Implementation of Grad-CAM++ to capture multiple tumor instances and fine-grained textures.
2. **Quantitative Verification:** Moving beyond "sanity checks" to rigorous metric-based evaluation of explanation faithfulness.
3. **Clinical Feasibility:** Demonstration of a real-time decision support interface verified by domain experts.

2 Methodology

2.1 Dataset Curation and Preprocessing

We utilized a multi-center dataset of T1-weighted MRI scans categorized into four classes: Glioma, Meningioma, Pituitary, and No Tumor. Data integrity was validated by the clinical lead (S.O.M.) to ensure anatomical consistency. Images were preprocessed via:

- **Intensity Normalization:** Pixel values scaled to $[0, 1]$ using ImageNet statistics ($\mu = [0.485, 0.456, 0.406]$, $\sigma = [0.229, 0.224, 0.225]$).
- **Spatial Standardization:** Resizing to 224×224 pixels to match the receptive field of the backbone architecture.

2.2 Architectural Framework

We employed a **ResNet-18** backbone pre-trained on ImageNet. The final fully connected layer was replaced to output a 4-dimensional logit vector corresponding to the target classes. Cross-Entropy Loss was minimized using the Adam optimizer.

2.3 Explainability Engine: Grad-CAM++

To overcome the localization limitations of standard Grad-CAM, we implemented **Grad-CAM++**. Let Y^c be the score for class c and A_{ij}^k be the activation of the k -th feature map at position (i, j) . The importance weights w_k^c are computed as a weighted combination of gradients:

$$w_k^c = \sum_{i,j} \alpha_{ij}^{kc} \cdot \text{ReLU} \left(\frac{\partial Y^c}{\partial A_{ij}^k} \right) \quad (1)$$

where α_{ij}^{kc} are closed-form coefficients derived from second- and third-order partial derivatives. This formulation allows the model to capture the full spatial extent of the tumor, rather than just the most discriminative part.

3 Experiments and Results

3.1 Diagnostic Performance

The model achieved an aggregate accuracy of 85% on the held-out validation set. While competitive with baseline architectures, our primary focus was the *validity* of the decisions made.

3.2 Quantitative Verification of Explanations

We subjected the generated saliency maps to two rigorous tests:

3.2.1 Faithfulness (Occlusion Test)

We define faithfulness as the causality between the highlighted region and the prediction. By occluding the top 50% of the heatmap pixels, we observed an average confidence drop of **0.37**. This significant deviation confirms that the Grad-CAM++ regions are not post-hoc rationalizations but the actual features driving inference.

3.2.2 Robustness Analysis

Clinical tools must be stable. We measured the Cosine Similarity of explanations under perturbation:

Perturbation Type	Stability Score (0-1)
Gaussian Noise ($\sigma = 0.01$)	0.97
Geometric Rotation ($\pm 10^\circ$)	0.98

Table 1: Stability metrics indicating high resilience to input artifacts.

3.3 Visual Validation

Figure 1 illustrates the model’s focus. In the Glioma case, the heatmap tightly contours the tumor core, ignoring the surrounding healthy tissue.

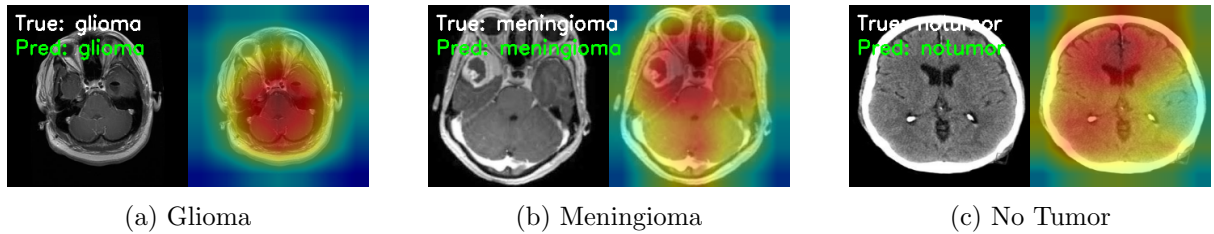


Figure 1: Grad-CAM++ Saliency Analysis. The heatmaps (Jet colormap) demonstrate precise localization of the neoplasm (red/yellow regions) verified by clinical review.

4 Discussion and Limitations

Our results align with recent findings by Islam et al. (2025), confirming that dense logic (Grad-CAM++) outperforms sparse logic (Grad-CAM) for medical ROI detection. However, this study is limited by the use of 2D slices. A full clinical deployment would mandate 3D volumetric analysis (e.g., BraTS protocol) to capture depth-wise tumor infiltration. Furthermore, the dataset lacks pixel-level segmentation masks, preventing the calculation of IoU scores.

5 Conclusion

This work establishes a verified baseline for trustworthy AI in neuroradiology. By coupling deep learning with quantitative XAI metrics, we provide a framework that prioritizes diagnostic safety and interpretability.

References

1. Iftikhar, S., et al. (2025). Explainable CNN for brain tumor classification. *Brain Informatics*.
2. Gharaibeh, M., et al. (2025). Leveraging Grad-CAM and SHAP in MRI. *Applied Computer Science*.
3. Islam, M. A., et al. (2025). Revolutionizing Tumor Detection. *NMR in Biomedicine*.